

# EVALUATING MODEL FIT

*Sri Kanajan*

---

## EVALUATING MODEL FIT

---

# LEARNING OBJECTIVES

- Define regularization, bias, and error metrics for regression problems
- Evaluate model fit using loss functions
- Select regression methods based on fit and complexity

---

**COURSE**

---

**PRE-WORK**

---

# PRE-WORK REVIEW

---

- Understand goodness of fit (r-squared)
- Measure statistical significance of feature coefficients
- Recall what a residual is
- Implement a sklearn estimator to predict a target variable

---

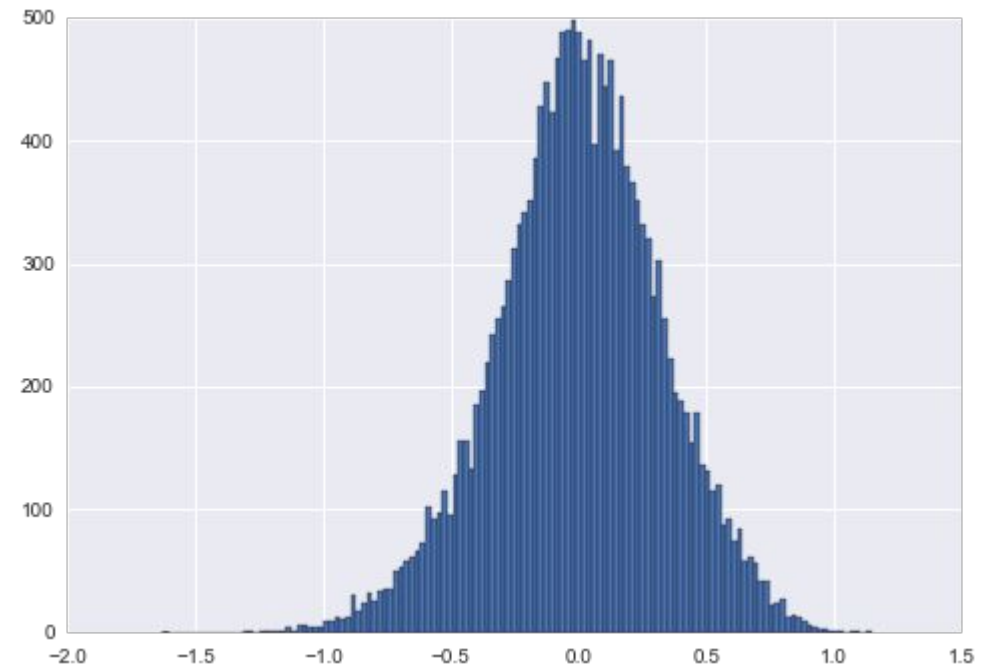
## INTRODUCTION

---

# LINEAR MODELS AND ERROR

## RECALL: WHAT'S RESIDUAL ERROR?

- ▶ In linear models, residual error must be normal with a median close to zero.
- ▶ Individual residuals are useful to see the error of specific points, but it doesn't provide an overall picture for optimization.
- ▶ We need a metric to summarize the error in our model into one value.
- ▶ Mean square error: the mean residual error in our model



---

## MEAN SQUARED ERROR (MSE)

---

- To calculate MSE:
  - Calculate the difference between each target  $y$  and the model's predicted value  $\hat{y}$  (i.e. the residual)
  - Square each residual.
  - Take the mean of the squared residual errors.

$$\text{MSE} = \frac{1}{n} \sum_{i=1}^n (y_i - \tilde{y}_i)^2$$

---

## MEAN SQUARED ERROR (MSE)

---

- sklearn's metrics module includes a mean\_squared\_error function.

```
from sklearn import metrics  
metrics.mean_squared_error(y, model.predict(X))
```



---

## MEAN SQUARED ERROR (MSE)

---

- For example, two arrays of the same values would have an MSE of 0.

```
from sklearn import metrics  
metrics.mean_squared_error([1, 2, 3, 4, 5], [1, 2, 3, 4, 5])  
0.0
```

# MEAN SQUARED ERROR (MSE)

---

- Two arrays with different values would have a positive MSE.

```
from sklearn import metrics
metrics.mean_squared_error([1, 2, 3, 4, 5], [5, 4, 3, 2, 1])
# (4^2 + 2^2 + 0^2 + 2^2 + 4^2) / 5
8.0
```

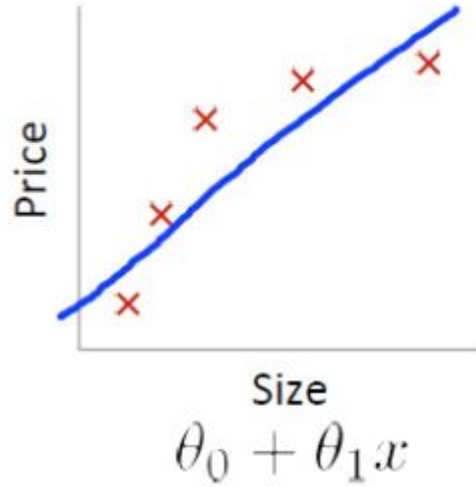
---

## HOW DO WE MINIMIZE ERROR?

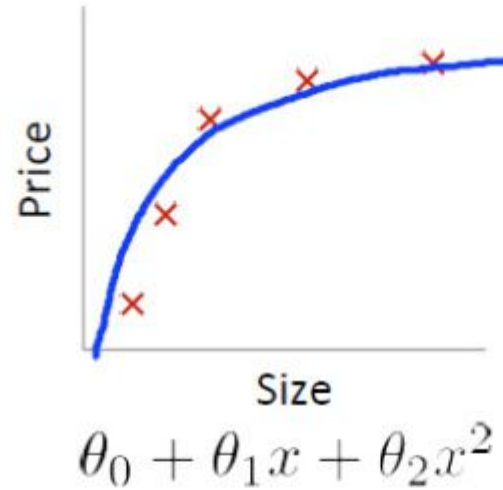
---

- The regression method we've used is called “Ordinary Least Squares”.
- This means that given a matrix  $X$ , solve for the *least* amount of square error for  $y$ .
- However, this assumes that  $X$  is unbiased, that it is representative of the population.
  - Otherwise, you are fitting the bias itself

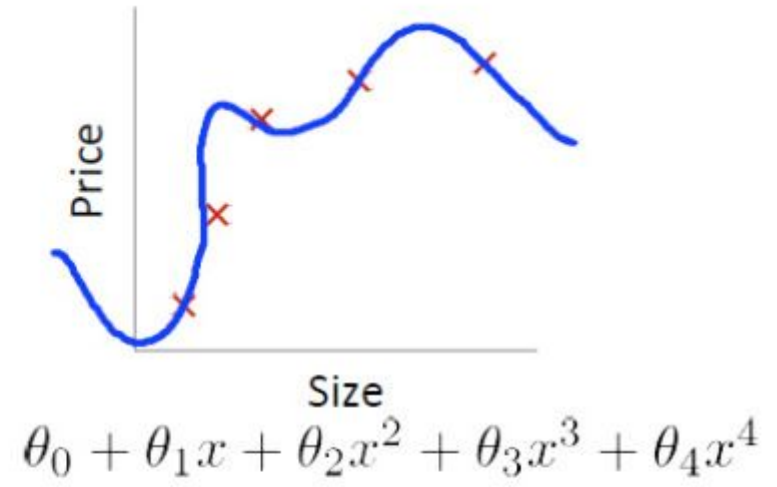
# BIAS VS. VARIANCE



High bias  
(underfit)  
 $d=1$



"Just right"  
 $d=2$



High variance  
(overfit)  
 $d=4$

**DEMO**

---

# CROSS VALIDATION

---

# CROSS VALIDATION

---

- Cross validation can help account for bias.
- The general idea is to
  - Generate several models on different cross sections of the data
  - Measure the performance of each
  - Take the mean performance
- This technique swaps bias error for generalized error, describing previous trends accurately enough to extend to future trends.

---

# K-FOLD CROSS VALIDATION

---

- k-fold cross validation
  - Split the data into  $k$  group
  - Train the model on all segments except one
  - Test model performance on the remaining set
- If  $k = 5$ , split the data into five segments and generate five models.

---

# USING K-FOLD CROSS VALIDATION WITH MSE

---

- Import the appropriate packages and load data.

```
from sklearn import cross_validation
wd = '../..../datasets/'
bikeshare = pd.read_csv(wd + 'bikeshare/bikeshare.csv')
weather = pd.get_dummies(bikeshare.weathersit, prefix='weather')
modeldata = bikeshare[['temp', 'hum']].join(weather[['weather_1',
'weather_2', 'weather_3']])
y = bikeshare.casual
```



---

## USING K-FOLD CROSS VALIDATION WITH MSE

---

- Build models on subsets of the data and calculate the average score.

```
kf = cross_validation.KFold(len(modeldata), n_folds=5, shuffle=True)
scores = []
for train_index, test_index in kf:
    lm =
linear_model.LinearRegression().fit(modeldata.iloc[train_index],
y.iloc[train_index])
    scores.append(metrics.mean_squared_error(y.iloc[test_index],
lm.predict(modeldata.iloc[test_index])))

print np.mean(scores)
```

---

## USING K-FOLD CROSS VALIDATION WITH MSE

---

- This can be compared to the model built on all of the data.
  - This score will be lower, but we're trading off bias error for generalized error:

```
lm = linear_model.LinearRegression().fit(modeldata, y)
print metrics.mean_squared_error(y, lm.predict(modeldata))
```
- Which approach would predict new data more accurately?

## **GUIDED PRACTICE**

---

# **CROSS VALIDATION WITH LINEAR REGRESSION**

---

# ACTIVITY: CROSS VALIDATION WITH LINEAR REGRESSION

---



## EXERCISE

### DIRECTIONS (20 minutes)

If we were to continue increasing the number of folds in cross validation, would error increase or decrease?

1. Using the previous code example, perform k-fold cross validation for all even numbers between 2 and 50.
2. Answer the following questions:
  - a. What does `shuffle=True` do?
  - b. At what point does cross validation no longer seem to help the model?
3. Hint: `range(2, 51, 2)` produces a list of even numbers from 2 to 50

### DELIVERABLE

Answers to questions

## **INTRODUCTION**

---

# **REGULARIZATION AND CROSS VALIDATION**

---

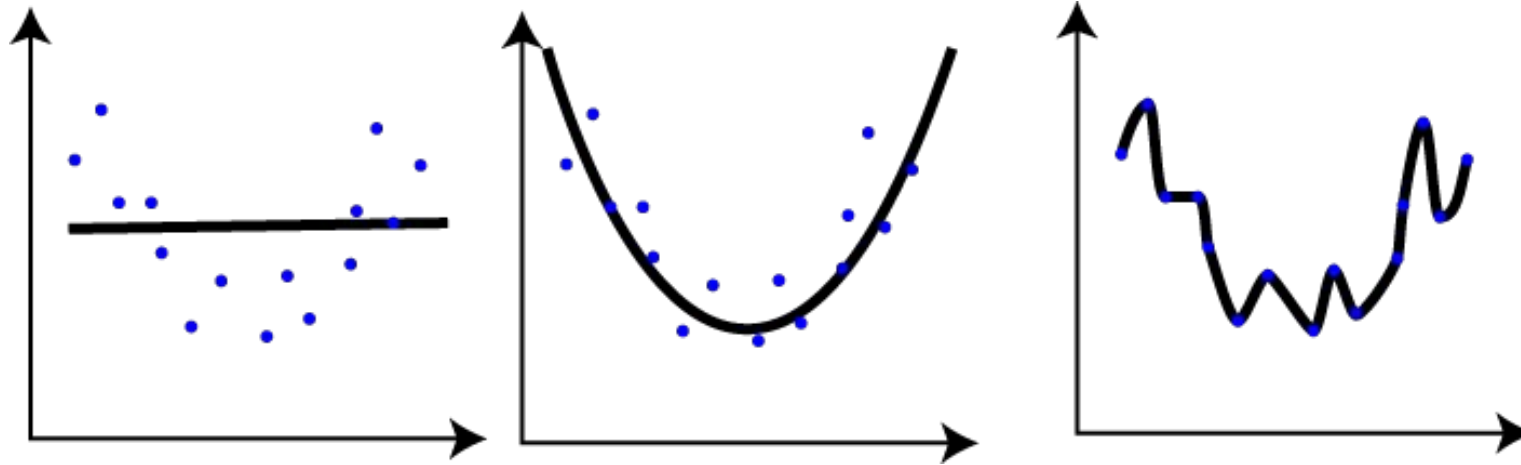
# WHAT IS REGULARIZATION? AND WHY DO WE USE IT?

---

- Regularization is an additive approach to protect models against overfitting (being potentially biased and overconfident, not generalizing well).
- Regularization becomes an additional weight to coefficients, shrinking them closer to zero.
- L1 (Lasso Regression) adds the extra weight to the sum of the absolute of the coefficients.
- L2 (Ridge Regression) adds the weight to the sum of the square of the coefficients.
- 
- Use Lasso when you want less features (sparse model) and Ridge when you want the most  $R^2$  but are okay with more features. Lasso will also drop features when they are somewhat correlated to each other

# WHAT IS OVERFITTING?

---



- The first model poorly explains the data.
- The second model explains the general curve of the data.
- The third model drastically overfits the model, bending to every point.
- Regularization helps prevent the third model.

---

## WHERE REGULARIZATION MAKES SENSE

---

- What happens to MSE if use Lasso or Ridge Regression directly?

```
lm = linear_model.LinearRegression().fit(modeldata, y)
print metrics.mean_squared_error(y, lm.predict(modeldata))
lm = linear_model.Lasso().fit(modeldata, y)
print metrics.mean_squared_error(y, lm.predict(modeldata))
lm = linear_model.Ridge().fit(modeldata, y)
print metrics.mean_squared_error(y, lm.predict(modeldata))
```

```
1672.58110765 # OLS
```

```
1725.41581608 # L1
```

```
1672.60490113 # L2
```



---

## WHERE REGULARIZATION MAKES SENSE

---

- It doesn't seem to help. Why is that?
- We need to optimize the regularization weight parameter (called  $\alpha$ ) through cross validation.

---

# ACTIVITY: KNOWLEDGE CHECK

---



## EXERCISE

### ANSWER THE FOLLOWING QUESTIONS (5 minutes)

1. Why is regularization important?
2. What does it protect against and how?

### DELIVERABLE

Answers to the above questions

**DEMO**

---

# UNDERSTANDING REGULARIZATION EFFECTS

---

## QUICK CHECK

---

- We are working with the bikeshare data to predict riders over hours/days with a few features.
- Does it make sense to use a ridge regression or a lasso regression?
- Why?

---

# UNDERSTANDING REGULARIZATION EFFECTS

---

- Let's test a variety of alpha weights for Ridge Regression on the bikeshare data.

```
alphas = np.logspace(-10, 10, 21)
for a in alphas:
    print 'Alpha:', a
    lm = linear_model.Ridge(alpha=a)
    lm.fit(modeldata, y)
    print lm.coef_
    print metrics.mean_squared_error(y, lm.predict(modeldata))
```

- What happens to the weights of the coefficients as alpha increases? What happens to the error as alpha increases?

---

## WE CAN MAKE THIS EASIER WITH GRID SEARCH!

---

- Grid search exhaustively searches through all given options to find the best solution. Grid search will try all combos given in `param_grid`.

```
param_grid = {  
    'intercept': [True, False],  
    'alpha': [1, 2, 3],  
}
```

---

## WE CAN MAKE THIS EASIER WITH GRID SEARCH!

---

▸ This param grid has six different options:

- intercept True, alpha 1
- intercept True, alpha 2
- intercept True, alpha 3
- intercept False, alpha 1
- intercept False, alpha 2
- intercept False, alpha 3

```
param_grid = {  
    'intercept': [True, False],  
    'alpha': [1, 2, 3],  
}
```

---

## WE CAN MAKE THIS EASIER WITH GRID SEARCH!

---

- This is an incredibly powerful, automated machine learning tool!

```
from sklearn import grid_search
```

```
alphas = np.logspace(-10, 10, 21)  
gs = grid_search.GridSearchCV(  
    estimator=linear_model.Ridge(),  
    param_grid={'alpha': alphas},  
    scoring='mean_squared_error')
```



---

## WE CAN MAKE THIS EASIER WITH GRID SEARCH!

---

```
gs.fit(modeldata, y)
```

```
print -gs.best_score_ # mean squared error here comes in negative, so  
let's make it positive.
```

```
print gs.best_estimator_ # explains which grid_search setup worked  
best
```

```
print gs.grid_scores_ # shows all the grid pairings and their  
performances.
```

## GUIDED PRACTICE

---

# GRID SEARCH CV, SOLVING FOR ALPHA

---

# ACTIVITY: GRID SEARCH CV, SOLVING FOR ALPHA

---



## EXERCISE

### DIRECTIONS (25 minutes)

1. Modify the previous code to do the following:
  - a. Introduce cross validation into the grid search. This is accessible from the cv argument.
  - b. Add `fit_intercept = True` and `False` to the `param_grid` dictionary.
  - c. Re-investigate the best score, best estimator, and grid score attributes as a result of the grid search.

### DELIVERABLE

New code and output that meets above requirements

## INTRODUCTION

---

# MINIMIZING LOSS THROUGH GRADIENT DESCENT

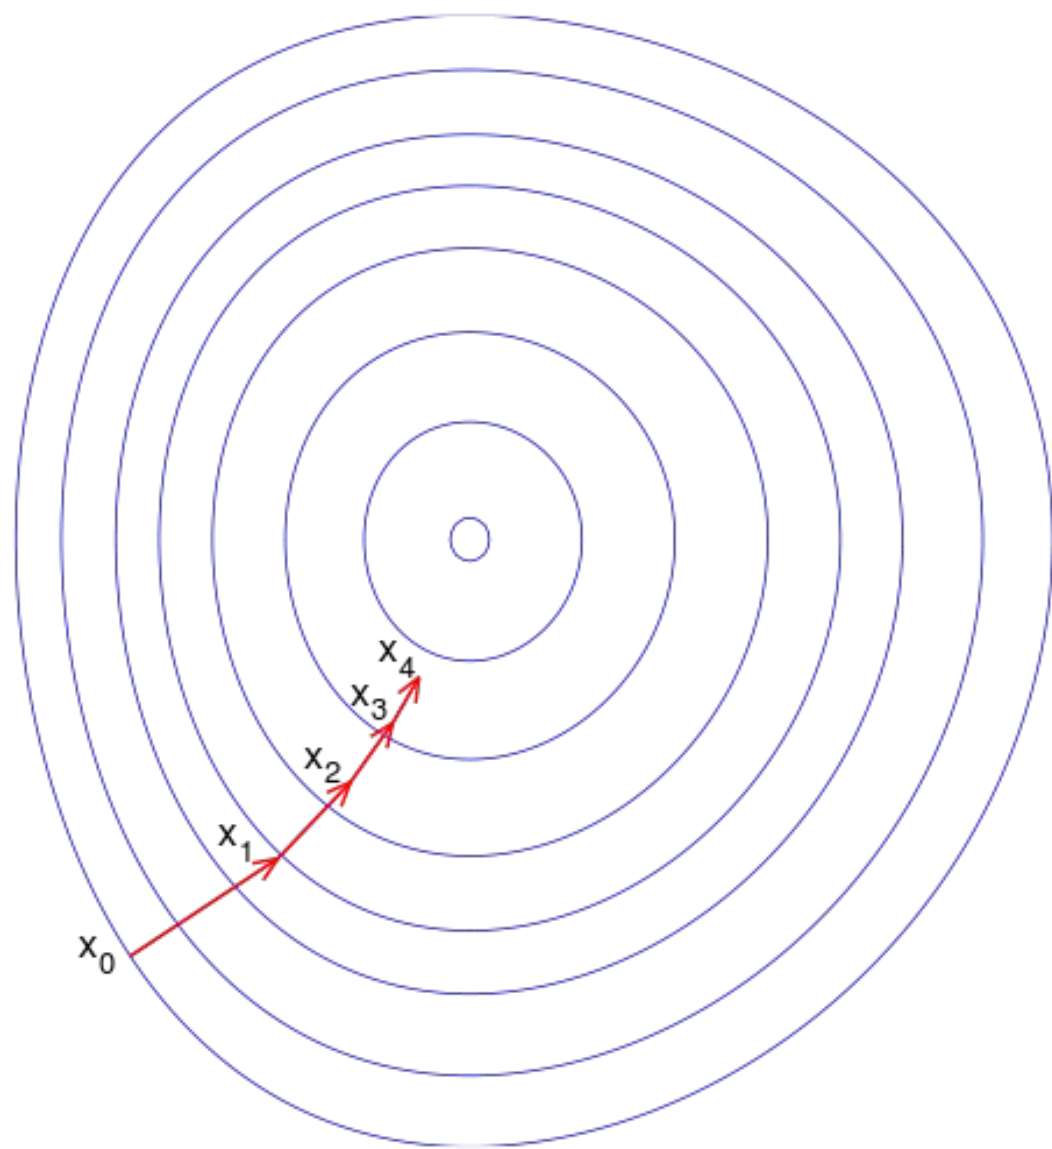
---

# GRADIENT DESCENT

---

- Gradient Descent can also help us minimize error.
- How Gradient Descent works:
  - A random linear solution is provided as a starting point
  - The solver attempts to find a next “step”: take a step in any direction and measure the performance.
  - If the solver finds a better solution (i.e. lower MSE), this is the new starting point.
  - Repeat these steps until the performance is optimized and no “next steps” perform better. The size of steps will shrink over time.

# GRADIENT DESCENT



# GLOBAL VS LOCAL MINIMUMS

---

- Gradient Descent could solve for a *local* minimum instead of a *global* minimum.
- A *local* minimum is confined to a very specific subset of solutions. The *global* minimum considers all solutions. These could be equal, but that's not always true.



---

**INDEPENDENT PRACTICE**

---

**ON YOUR OWN**



---

# ACTIVITY: ON YOUR OWN

---



## EXERCISE

### DIRECTIONS (30 minutes)

There are tons of ways to approach a regression problem.

1. Demonstrate the `grid_search` module.
2. Use a model you evaluated last class or the simpler one from today. Implement `param_grid` in grid search to answer the following questions:
  - a. With a set of values between  $10^{-10}$  and  $10^{-1}$ , how does MSE change?
  - b. Our data suggests we use L1 regularization. Using a grid search with `l1_ratios` between 0 and 1, increasing every 0.05, does this statement hold true?

### DELIVERABLE

Answered questions

# ACTIVITY: ON YOUR OWN



## EXERCISE

### Starter Code

```
params = {} # put your gradient descent parameters here
gs = grid_search.GridSearchCV(
    estimator=linear_model.SGDRegressor(),
    cv=cross_validation.KFold(len(modeldata), n_folds=5, shuffle=True),
    param_grid=params,
    scoring='mean_squared_error',
)

gs.fit(modeldata, y)

print 'BEST ESTIMATOR'
print -gs.best_score_
print gs.best_estimator_
print 'ALL ESTIMATORS'
print gs.grid_scores_
```

---

**CONCLUSION**

---

# TOPIC REVIEW

---

## LESSON REVIEW

---

- What's the (typical) range of r-squared?
- What's the range of mean squared error?
- How would changing the scale or interpretation of  $y$  (your target variable) effect mean squared error?
- What's cross validation, and why do we use it in machine learning?
- What is error due to bias? What is error due to variance? Which is better for a model to have, if it had to have one?
- How does gradient descent try a different approach to minimizing error?

---

## **BEFORE NEXT CLASS**

---

- Project: Final Project, Deliverable 1

---

**LESSON**

---

**Q & A**

## **LESSON**

---

# **EXIT TICKET**

**DON'T FORGET TO FILL OUT YOUR EXIT TICKET**